

Detecting Applications in Traffic Flows Using Machine Learning

Authors: Sam Xie (muyanxiesam), Jillian Chen (jianghui)

Project Overview

With the increasing use of encryption in network traffic, traditional methods of application identification based on payload inspection have become less effective. This project aims to explore machine learning techniques to identify applications based on **traffic flow characteristics**, such as packet size distributions, timing patterns, and flow duration. Accurate application identification is crucial for network management, security monitoring, and quality of service enforcement.

Flow Definition

A **network flow** is defined as a sequence of packets sharing the same 5-tuple:

- Source IP address
- Destination IP address
- Source port
- Destination port
- Protocol (TCP/UDP)

This project extracts flows from raw PCAP data and computes statistical features for each flow, enabling flow-level application classification.

This project will aim to reproduce results from the [netml vpn non-vpn dataset paper](#)

Data

The project will utilize a dataset of network traffic captures in PCAPNG format, specifically focusing on encrypted traffic. The dataset will be processed to:

1. Extract individual packets with their metadata
2. Group packets into flows using the 5-tuple definition
3. Compute statistical flow features (packet counts, size distributions, timing statistics)
4. Aggregate packet-level labels to flow-level labels

A link to the dataset is here: [Dataset Link](#).

Data Processing

checking the required datasets

```
[1]: # detect if we have all the packages and data we need
from scapy.all import rdpcap, wrpcap
from pathlib import Path
import os

def fix_pcap_timestamps(input_file, output_file):
    """Remove first packet and normalize timestamps."""

```

```

print(f"Loading PCAP file...")
packets = rdpcap(str(input_file))
print(f"Original packet count: {len(packets)}")

# Remove first packet
packets_fixed = packets[1:]

# Set timestamps to valid absolute values starting from a base time
# Use a recent but fixed epoch time (e.g., Jan 1, 2020)
base_time = 1577836800.0 # 2020-01-01 00:00:00 UTC

if len(packets_fixed) > 0:
    print(f"Normalizing timestamps to start from 2020-01-01...")
    for i, pkt in enumerate(packets_fixed):
        # Set timestamp as base_time + packet index (1 second apart)
        # This ensures all timestamps are valid and sequential
        pkt.time = base_time + (i * 0.001) # 1ms apart

print(f"Fixed packet count: {len(packets_fixed)}")
print(f"Writing to {output_file}...")
wrpcap(str(output_file), packets_fixed)
print(f"\u2713 Saved successfully!")

# File paths
original_pcap = Path("data/traffic.pcapng")
fixed_pcap = Path("data/traffic_fixed.pcap")
extracted_csv = Path("data/extracted.csv")

# Fix PCAP timestamps if needed
if not fixed_pcap.exists():
    print(f"Fixing timestamps in {original_pcap}...")
    if not original_pcap.exists():
        raise FileNotFoundError(f"Required file {original_pcap} not found.")
    fix_pcap_timestamps(original_pcap, fixed_pcap)
    print("\u2713 Timestamps fixed successfully!")
else:
    print(f"\u2713 Using existing fixed PCAP: {fixed_pcap}")

# Check if extracted.csv exists
if not extracted_csv.exists():
    print(f"\n⚠ Warning: {extracted_csv} not found.")
    print(f"Extracting comments from original PCAP...")
    os.system("tshark -r data/traffic.pcapng -T fields -E header=y -e frame.number")
    print("\u2713 Comments extracted.")
else:
    print(f"\u2713 Found {extracted_csv}")

print(f"\n\u2713 All required files are ready!")

```

WARNING: No IPv4 address found on anpi3 !
 WARNING: No IPv4 address found on anpi0 !
 WARNING: No IPv4 address found on anpi0 !
 WARNING: more No IPv4 address found on anpi1 !
 WARNING: more No IPv4 address found on anpi1 !
 ✓ Using existing fixed PCAP: data/traffic_fixed.pcap
 ✓ Found data/extracted.csv

✓ All required files are ready!

Flow Extraction and Feature Engineering

We extract flows from the PCAP file using the 5-tuple definition (src_ip, dst_ip, src_port, dst_port, protocol). For each flow, we compute statistical features that capture the behavior of the traffic without requiring payload inspection.

```
[2]: import pandas as pd
import numpy as np
from scapy.all import rdpcap, IP, TCP, UDP, IPv6

# Load packet-level observations (labels)
observations = pd.read_csv('data/extracted.csv',
                           skiprows=2,
                           names=['id', 'frame_number', 'frame_comment'])
observations['id'] = observations['id'] - 2 # Adjust for removed first packet

print(f"Loaded {len(observations)} packet labels")

# Load packets from PCAP
fixed_pcap = 'data/traffic_fixed.pcap'
print(f"\nLoading packets from {fixed_pcap}...")
packets = rdpcap(fixed_pcap)
print(f"✓ Loaded {len(packets)} packets")

# Extract packet information with flow keys
print("\nExtracting packet information and flow keys...")
packet_data = []

for idx, pkt in enumerate(packets):
    # Extract flow 5-tuple
    src_ip = dst_ip = src_port = dst_port = protocol = None
    pkt_len = len(pkt)
    timestamp = float(pkt.time)

    # Handle IPv4
    if IP in pkt:
        src_ip = pkt[IP].src
        dst_ip = pkt[IP].dst
        protocol = pkt[IP].proto
    # Handle IPv6
    elif IPv6 in pkt:
        src_ip = pkt[IPv6].src
        dst_ip = pkt[IPv6].dst
        protocol = pkt[IPv6].nh

    # Extract ports for TCP/UDP
    if TCP in pkt:
        src_port = pkt[TCP].sport
        dst_port = pkt[TCP].dport
        protocol = 6 # TCP
    elif UDP in pkt:
        src_port = pkt[UDP].sport
        dst_port = pkt[UDP].dport
        protocol = 17 # UDP

    # Create bidirectional flow key (sorted to make it direction-agnostic)
    if src_ip and dst_ip:
        endpoints = sorted([(src_ip, src_port or 0), (dst_ip, dst_port or 0)])
        flow_key = f"{endpoints[0][0]}:{endpoints[0][1]}-{endpoints[1][0]}:{endpoints[1][1]}"
    else:
        flow_key = f"unknown_{idx}"

    packet_data.append({
```

```

    'packet_idx': idx,
    'flow_key': flow_key,
    'src_ip': src_ip,
    'dst_ip': dst_ip,
    'src_port': src_port,
    'dst_port': dst_port,
    'protocol': protocol,
    'pkt_len': pkt_len,
    'timestamp': timestamp
})

# Create packet DataFrame
pdf = pd.DataFrame(packet_data)
print(f"\u2713 Extracted {len(pdf)} packet records")
print(f"\u2713 Found {pdf['flow_key'].nunique()} unique flows")

```

Loaded 529018 packet labels

Loading packets from data/traffic_fixed.pcap...
 ✓ Loaded 529018 packets

Extracting packet information and flow keys...
 ✓ Loaded 529018 packets

Extracting packet information and flow keys...
 ✓ Extracted 529018 packet records
 ✓ Found 111284 unique flows
 ✓ Extracted 529018 packet records
 ✓ Found 111284 unique flows

[3]: # Merge packet data with labels

```

pdf = pdf.merge(observations[['id', 'frame_comment']],
                left_on='packet_idx', right_on='id', how='left')
pdf = pdf.drop(columns=['id'])
pdf = pdf.rename(columns={'frame_comment': 'label'})

print(f"Packets with labels: {pdf['label'].notna().sum()}")
print(f"\nSample packet data:")
pdf.head()

```

Packets with labels: 529018

Sample packet data:

[3]:

	packet_idx	flow_key	src_ip	dst_ip	src_port	dst_port	protocol	pkt_len	timestamp
0	0	131.202.240.87:8118-222.186.31.178:49445-6	222.186.31.178	131.202.240.87	49445	8118	6	60	1.577837e+
1	1	131.202.240.87:4506-207.241.228.202:80-6	131.202.240.87	207.241.228.202	4506	80	6	66	1.577837e+
2	2	131.202.240.87:4506-207.241.228.202:80-6	207.241.228.202	131.202.240.87	80	4506	6	66	1.577837e+
3	3	131.202.240.87:4506-207.241.228.202:80-6	131.202.240.87	207.241.228.202	4506	80	6	54	1.577837e+
4	4	131.202.240.87:4506-207.241.228.202:80-6	131.202.240.87	207.241.228.202	4506	80	6	240	1.577837e+

[4]: # Check flow distribution

```

print("Flow size distribution:")
flow_sizes = pdf.groupby('flow_key').size()
print(f" Min packets per flow: {flow_sizes.min()}")
print(f" Max packets per flow: {flow_sizes.max()}")
print(f" Mean packets per flow: {flow_sizes.mean():.2f}")
print(f" Median packets per flow: {flow_sizes.median():.2f}")

```

```
print(f"\nLabel distribution in packets:")
print(pdf['label'].value_counts().head(10))
```

```
Flow size distribution:
Min packets per flow: 1
Max packets per flow: 875
Mean packets per flow: 4.75
Median packets per flow: 2.00
```

```
Label distribution in packets:
```

```
label
audio_hangouts_hangouts-audio    126174
audio_facebook_facebook-audio   122003
file-transfer_skype_skype-file   85035
audio_skype_skype-audio        58558
email_email_email                18876
p2p_torrent_torrent              14926
video_skype_skype-video          13539
chat_skype_skype-chat            12542
video_youtube_youtube            11571
audio_voipbuster_voipbuster     11355
Name: count, dtype: int64
label
audio_hangouts_hangouts-audio    126174
audio_facebook_facebook-audio   122003
file-transfer_skype_skype-file   85035
audio_skype_skype-audio        58558
email_email_email                18876
p2p_torrent_torrent              14926
video_skype_skype-video          13539
chat_skype_skype-chat            12542
video_youtube_youtube            11571
audio_voipbuster_voipbuster     11355
Name: count, dtype: int64
```

Flow Feature Engineering

For each flow, we compute statistical features that capture traffic patterns without inspecting payloads:

Packet Count Features:

- Total packets in flow
- Forward/backward packet counts

Packet Size Features:

- Total bytes, mean, std, min, max packet sizes
- Forward/backward size statistics

Timing Features:

- Flow duration
- Inter-arrival time statistics (mean, std, min, max)

Protocol Features:

- Protocol type (TCP/UDP)
- Port numbers

```
[5]: def extract_flow_features(flow_df):
    """Extract statistical features for a single flow."""
    features = {}
```

```

# Basic flow info
features['flow_key'] = flow_df['flow_key'].iloc[0]

# Packet count features
features['total_packets'] = len(flow_df)
features['total_bytes'] = flow_df['pkt_len'].sum()

# Packet size statistics
features['pkt_len_mean'] = flow_df['pkt_len'].mean()
features['pkt_len_std'] = flow_df['pkt_len'].std() if len(flow_df) > 1 else 0
features['pkt_len_min'] = flow_df['pkt_len'].min()
features['pkt_len_max'] = flow_df['pkt_len'].max()
features['pkt_len_median'] = flow_df['pkt_len'].median()

# Timing features
timestamps = flow_df['timestamp'].sort_values()
features['flow_duration'] = timestamps.max() - timestamps.min()

if len(timestamps) > 1:
    iat = timestamps.diff().dropna() # Inter-arrival times
    features['iat_mean'] = iat.mean()
    features['iat_std'] = iat.std()
    features['iat_min'] = iat.min()
    features['iat_max'] = iat.max()
else:
    features['iat_mean'] = 0
    features['iat_std'] = 0
    features['iat_min'] = 0
    features['iat_max'] = 0

# Protocol features
features['protocol'] = flow_df['protocol'].iloc[0] if flow_df['protocol'].notna() else 0

# Port features (use min port as it's often the server port)
src_ports = flow_df['src_port'].dropna()
dst_ports = flow_df['dst_port'].dropna()
all_ports = pd.concat([src_ports, dst_ports])
features['min_port'] = all_ports.min() if len(all_ports) > 0 else 0
features['max_port'] = all_ports.max() if len(all_ports) > 0 else 0

# Packet size distribution features
features['small_packets'] = (flow_df['pkt_len'] < 100).sum() # Small packets
features['medium_packets'] = ((flow_df['pkt_len'] >= 100) & (flow_df['pkt_len'] < 500)).sum()
features['large_packets'] = (flow_df['pkt_len'] >= 500).sum()

# Bytes per second (flow rate)
if features['flow_duration'] > 0:
    features['bytes_per_sec'] = features['total_bytes'] / features['flow_duration']
    features['packets_per_sec'] = features['total_packets'] / features['flow_duration']
else:
    features['bytes_per_sec'] = features['total_bytes']
    features['packets_per_sec'] = features['total_packets']

# Get the most common label for this flow (majority voting)
labels = flow_df['label'].dropna()
if len(labels) > 0:
    features['label'] = labels.mode().iloc[0]
else:
    features['label'] = None

return features

```

```

print("Flow feature extraction function defined ✓")
Flow feature extraction function defined ✓

[6]: # Extract features for all flows
print("Extracting flow features...")
from tqdm import tqdm

flow_features = []
grouped = pdf.groupby('flow_key')
total_flows = len(grouped)

for flow_key, flow_df in tqdm(grouped, total=total_flows, desc="Processing flows"):
    features = extract_flow_features(flow_df)
    flow_features.append(features)

# Create flow DataFrame
flow_df = pd.DataFrame(flow_features)

# Remove flows without labels
flow_df = flow_df.dropna(subset=['label'])

print(f"\n✓ Extracted {len(flow_df)} labeled flows")
print(f"✓ {len(flow_df.columns) - 2} features per flow") # -2 for flow_key and label
print("Flow features:")
print(list(flow_df.columns))

```

Extracting flow features...

Processing flows: 100%|██████████| 111284/111284 [01:04<00:00, 1714.58it/s]

✓ Extracted 111284 labeled flows
✓ 20 features per flow

Flow features:
['flow_key', 'total_packets', 'total_bytes', 'pkt_len_mean', 'pkt_len_std', 'pkt_le

Train and Test Split

We split the flow dataset into training and testing sets (80/20 split), ensuring that the split maintains the distribution of application classes through stratification.

```

[7]: from sklearn.model_selection import train_test_split

# Preview the flow data
print("Flow data summary:")
print(flow_df.describe())
print("Label distribution:")
print(flow_df['label'].value_counts())

# Separate features and labels
feature_cols = [col for col in flow_df.columns if col not in ['flow_key', 'label']]
X = flow_df[feature_cols]
y = flow_df['label']

print("\n✓ Feature matrix shape: {X.shape}")
print("✓ Number of classes: {y.unique()}")

```

Flow data summary:

	total_packets	total_bytes	pkt_len_mean	pkt_len_std
count	111284	111284	111284	111284
mean	100.0	100.0	100.0	100.0
std	100.0	100.0	100.0	100.0
min	0.0	0.0	0.0	0.0
max	1000.0	1000.0	1000.0	1000.0
q1	50.0	50.0	50.0	50.0
q3	150.0	150.0	150.0	150.0

count	111284.000000	111284.000000	111284.000000	111284.000000	111284.000000
mean	4.753765	793.260972	96.330745	27.205372	
std	14.437535	5706.271733	77.464731	85.395472	
min	1.000000	54.000000	49.000000	0.000000	
25%	2.000000	128.000000	64.000000	0.000000	
50%	2.000000	256.000000	64.000000	0.000000	
75%	4.000000	384.000000	129.500000	0.000000	
max	875.000000	399301.000000	1196.460000	758.060903	
count	111284.000000	111284.000000	111284.000000	111284.000000	111284.000000
mean	76.211342	129.843715	93.802429	83.561406	
std	27.928165	198.008331	89.163239	140.379092	
min	44.000000	54.000000	44.000000	0.000000	
25%	64.000000	64.000000	64.000000	0.001000	
50%	64.000000	64.000000	64.000000	0.001000	
75%	73.000000	131.000000	110.000000	149.557750	
max	816.000000	1500.000000	1500.000000	512.397000	
count	111284.000000	5.625700e+04	111284.000000	111284.000000	111284.000000
mean	23.328593	6.441239e+01	5.511464	72.004099	
std	49.332916	7.0044665e+01	38.588282	120.810929	
min	0.000000	0.000000e+00	0.000000	0.000000	
25%	0.001000	1.101272e-07	0.001000	0.001000	
50%	0.001000	5.370946e+01	0.001000	0.001000	
75%	33.892500	1.102878e+02	0.001000	137.765000	
max	498.804000	3.581694e+02	498.804000	511.023000	
count	111284.000000	111284.000000	111284.000000	111284.000000	111284.000000
mean	16.588306	5981.056100	54292.204935	3.661119	
std	2.087866	5562.934703	11625.858899	9.788265	
min	6.000000	19.000000	68.000000	0.000000	
25%	17.000000	5355.000000	51991.000000	2.000000	
50%	17.000000	5355.000000	56597.000000	2.000000	
75%	17.000000	5355.000000	61009.000000	4.000000	
max	17.000000	61009.000000	65535.000000	863.000000	
count	111284.000000	111284.000000	bytes_per_sec	packets_per_sec	
mean	0.779474	0.313172	1.112840e+05	111284.000000	
std	7.768443	4.004202	1.373792e+05	1069.973603	
min	0.000000	0.000000	4.368990e-01	0.004010	
25%	0.000000	0.000000	2.374907e+00	0.035798	
50%	0.000000	0.000000	1.279788e+05	1333.324009	
75%	1.000000	0.000000	1.280093e+05	2000.144969	
max	504.000000	340.000000	1.632118e+06	2000.144969	
Label distribution:					
label					
audio_facebook_facebook-audio	38940				
audio_hangouts_hangouts-audio	33038				
file-transfer_skype_skype-file	14127				
audio_skype_skype-audio	10539				
audio_voipbuster_voipbuster	2660				
email_email_email	2114				
chat_skype_skype-chat	1893				
video_hangouts_hangouts-video	1469				
p2pTorrent_torrent	958				
video_youtube_youtube	826				
file-transfer_ftps_ftps-down	586				
video_skype_skype-video	512				
chat_facebook_facebook-chat	453				
chat_gmail_gmail-chat	429				
video_vimeo_vimeo	399				
video_facebook_facebook-video	388				
chat_icq_icq-chat	340				
chat_aim_aim-chat	339				

```

chat_hangouts_hangouts-chat      298
video_netflix_netflix           247
audio_spotify_spotify           198
file-transfer_ftps_ftps-up     134
tor_youtube_tor-youtube        87
file-transfer_scp_scp-down     84
file-transfer_scp_scp-up       82
file-transfer_sftp_sftp-down   72
file-transfer_sftp_sftp-up     63
video_vimeo_tor-vimeo          6
tor_facebook_tor-facebook      2
tor_twitter_tor-twitter        1
Name: count, dtype: int64

```

✓ Feature matrix shape: (111284, 20)
✓ Number of classes: 30

[8]: # Filter classes with too few samples for stratification

```

min_samples = 2
class_counts = y.value_counts()
valid_classes = class_counts[class_counts >= min_samples].index
mask = y.isin(valid_classes)
X_filtered = X[mask]
y_filtered = y[mask]

print(f"Classes with >= {min_samples} samples: {len(valid_classes)}")
print(f"Flows after filtering: {len(X_filtered)}")

# Perform 80/20 split with stratification
X_train, X_test, y_train, y_test = train_test_split(
    X_filtered,
    test_size=0.2,
    random_state=42,
    stratify=y_filtered
)

print(f"\nTraining set size: {len(X_train)} flows")
print(f"Test set size: {len(X_test)} flows")
print(f"\nTraining set class distribution:")
print(y_train.value_counts().head(10))
print(f"\nTest set class distribution:")
print(y_test.value_counts().head(10))

```

Classes with >= 2 samples: 29
Flows after filtering: 111283

Training set size: 89026 flows
Test set size: 22257 flows

Training set class distribution:

label	count
audio_facebook_facebook-audio	31152
audio_hangouts_hangouts-audio	26430
file-transfer_skype_skype-file	11302
audio_skype_skype-audio	8431
audio_voipbuster_voipbuster	2128
email_email_email	1691
chat_skype_skype-chat	1514
video_hangouts_hangouts-video	1175
p2p_torrent_torrent	766
video_youtube_youtube	661

Name: count, dtype: int64

Test set class distribution:

label	count
audio_facebook_facebook-audio	7788

```

audio_hangouts_hangouts-audio      6608
file-transfer_skype_skype-file    2825
audio_skype_skype-audio          2108
audio_voipbuster_voipbuster       532
email_email_email                  423
chat_skype_skype-chat              379
video_hangouts_hangouts-video     294
p2p_torrent_torrent                192
video_youtube_youtube              165
Name: count, dtype: int64

```

Model Training

We experiment with several machine learning algorithms using flow-level features:

1. **Random Forest** - Ensemble method, good for feature importance analysis
2. **Support Vector Machines (SVM)** - Effective for high-dimensional data
3. **Neural Networks (MLP)** - Can learn complex non-linear patterns

1. Random Forest

```
[9]: # Prepare flow features for training
print("Flow features being used:")
print(X_train.columns.tolist())
print(f"\nFeature data types:")
print(X_train.dtypes)

# Check for any missing values
print(f"\nMissing values per feature:")
print(X_train.isnull().sum())

# Fill any remaining NaN values with 0
X_train_processed = X_train.fillna(0)
X_test_processed = X_test.fillna(0)

print(f"\n Training features ready: {X_train_processed.shape}")
```

Flow features being used:
['total_packets', 'total_bytes', 'pkt_len_mean', 'pkt_len_std', 'pkt_len_min', 'pkt_len_max', 'pkt_len_median', 'flow_duration', 'iat_mean', 'iat_std', 'iat_min', 'iat_max', 'protocol', 'min_port', 'max_port', 'small_packets', 'medium_packets', 'large_packets', 'bytes_per_sec', 'packets_per_sec']

Feature data types:

Feature	Type
total_packets	int64
total_bytes	int64
pkt_len_mean	float64
pkt_len_std	float64
pkt_len_min	int64
pkt_len_max	int64
pkt_len_median	float64
flow_duration	float64
iat_mean	float64
iat_std	float64
iat_min	float64
iat_max	float64
protocol	int64
min_port	int64
max_port	int64
small_packets	int64
medium_packets	int64
large_packets	int64
bytes_per_sec	float64
packets_per_sec	float64

dtype: object

```
Missing values per feature:  
total_packets          0  
total_bytes            0  
pkt_len_mean           0  
pkt_len_std             0  
pkt_len_min             0  
pkt_len_max             0  
pkt_len_median          0  
flow_duration           0  
iat_mean                0  
iat_std                 43941  
iat_min                0  
iat_max                0  
protocol               0  
min_port                0  
max_port                0  
small_packets            0  
medium_packets           0  
large_packets            0  
bytes_per_sec            0  
packets_per_sec          0  
dtype: int64
```

✓ Training features ready: (89026, 20)

```
[10]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, classification_report  
import time  
  
# Initialize Random Forest Classifier  
rf_model = RandomForestClassifier(  
    n_estimators=100,          # Number of trees in the forest  
    max_depth=20,            # Maximum depth of trees  
    min_samples_split=5,     # Minimum samples required to split a node  
    min_samples_leaf=2,      # Minimum samples required at leaf node  
    random_state=42,         # For reproducibility  
    n_jobs=-1                # Use all available processors  
)  
  
print("Training Random Forest model on flow features...")  
start_time = time.time()  
rf_model.fit(X_train_processed, y_train)  
rf_train_time = time.time() - start_time  
print(f"\n✓ Training complete in {rf_train_time:.2f} seconds!")  
  
# Make predictions  
print("\nMaking predictions on test set...")  
y_pred_rf = rf_model.predict(X_test_processed)  
  
# Evaluate the model  
print("\n" + "="*60)  
print("RANDOM FOREST MODEL EVALUATION (Flow-Level)")  
print("=".*60)  
  
# Accuracy  
accuracy = accuracy_score(y_test, y_pred_rf)  
print(f"\nAccuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")  
  
# Calculate feature importance  
feature_importance = pd.DataFrame({  
    'feature': X_train_processed.columns,  
    'importance': rf_model.feature_importances_}  
).sort_values('importance', ascending=False)
```

```

print(f"\nTop 10 Most Important Flow Features:")
print(feature_importance.head(10))

Training Random Forest model on flow features...
✓ Training complete in 1.41 seconds!

Making predictions on test set...

=====
RANDOM FOREST MODEL EVALUATION (Flow-Level)
=====

Accuracy: 0.5039 (50.39%)

Top 10 Most Important Flow Features:
      feature  importance
14      max_port    0.175336
18    bytes_per_sec   0.103569
11        iat_max    0.092472
7     flow_duration   0.070209
19  packets_per_sec   0.062296
9        iat_std     0.057306
3      pkt_len_std    0.054664
6    pkt_len_median   0.052264
8        iat_mean    0.052191
1    total_bytes     0.048961
✓ Training complete in 1.41 seconds!

Making predictions on test set...

=====
RANDOM FOREST MODEL EVALUATION (Flow-Level)
=====

Accuracy: 0.5039 (50.39%)

Top 10 Most Important Flow Features:
      feature  importance
14      max_port    0.175336
18    bytes_per_sec   0.103569
11        iat_max    0.092472
7     flow_duration   0.070209
19  packets_per_sec   0.062296
9        iat_std     0.057306
3      pkt_len_std    0.054664
6    pkt_len_median   0.052264
8        iat_mean    0.052191
1    total_bytes     0.048961

```

2. Support Vector Machines (SVM)

```

[11]: from sklearn.svm import LinearSVC
from sklearn.preprocessing import StandardScaler
import warnings

warnings.filterwarnings('ignore')

# Scale features for SVM (SVMs are sensitive to feature scaling)
print("Scaling flow features for SVM...")
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_processed)
X_test_scaled = scaler.transform(X_test_processed)
print("✓ Scaling complete!")

# Initialize Linear SVM

```

```

svm_model = LinearSVC(
    C=1.0,                      # Regularization parameter
    max_iter=2000,                # Maximum iterations
    random_state=42,
    dual=True
)

print("\nTraining Linear SVM model on flow features...")
start_time = time.time()
svm_model.fit(X_train_scaled, y_train)
svm_train_time = time.time() - start_time
print(f"\n✓ Training complete in {svm_train_time:.2f} seconds!")

# Make predictions
print("\nMaking predictions on test set...")
y_pred_svm = svm_model.predict(X_test_scaled)

# Evaluate
print("\n" + "="*60)
print("SUPPORT VECTOR MACHINE (SVM) MODEL EVALUATION (Flow-Level)")
print("="*60)

accuracy_svm = accuracy_score(y_test, y_pred_svm)
print(f"\nAccuracy: {accuracy_svm:.4f} ({accuracy_svm*100:.2f}%)")

```

Scaling flow features for SVM...
✓ Scaling complete!

Training Linear SVM model on flow features...
✓ Training complete in 314.05 seconds!

Making predictions on test set...

=====
SUPPORT VECTOR MACHINE (SVM) MODEL EVALUATION (Flow-Level)
=====

Accuracy: 0.4198 (41.98%)
✓ Training complete in 314.05 seconds!

Making predictions on test set...

=====
SUPPORT VECTOR MACHINE (SVM) MODEL EVALUATION (Flow-Level)
=====

Accuracy: 0.4198 (41.98%)

3. Neural Networks

```

[12]: from sklearn.neural_network import MLPClassifier

# Initialize Multi-Layer Perceptron (Neural Network)
nn_model = MLPClassifier(
    hidden_layer_sizes=(128, 64, 32), # Three hidden layers
    activation='relu',
    solver='adam',                  # L2 regularization
    alpha=0.0001,
    batch_size=256,
    learning_rate='adaptive',
    learning_rate_init=0.001,
    max_iter=300,
    random_state=42,
)

```

```

    verbose=True,
    early_stopping=True,
    validation_fraction=0.1,
    n_iter_no_change=15
)

print("Training Neural Network model on flow features...")
print("Architecture: Input -> 128 -> 64 -> 32 -> Output\n")
start_time = time.time()
nn_model.fit(X_train_scaled, y_train)
nn_train_time = time.time() - start_time
print(f"\n Training complete in {nn_train_time:.2f} seconds!")
print(f"Training stopped at iteration: {nn_model.n_iter_}")

# Make predictions
y_pred_nn = nn_model.predict(X_test_scaled)

# Evaluate
print("\n" + "="*60)
print("NEURAL NETWORK MODEL EVALUATION (Flow-Level)")
print("="*60)

accuracy_nn = accuracy_score(y_test, y_pred_nn)
print(f"\nAccuracy: {accuracy_nn:.4f} ({accuracy_nn*100:.2f}%)")

```

Training Neural Network model on flow features...
Architecture: Input -> 128 -> 64 -> 32 -> Output

Iteration 1, loss = 1.72302430
Validation score: 0.437830
Iteration 2, loss = 1.49745900
Validation score: 0.438504
Iteration 1, loss = 1.72302430
Validation score: 0.437830
Iteration 2, loss = 1.49745900
Validation score: 0.438504
Iteration 3, loss = 1.46355719
Validation score: 0.435584
Iteration 4, loss = 1.44245657
Validation score: 0.447153
Iteration 3, loss = 1.46355719
Validation score: 0.435584
Iteration 4, loss = 1.44245657
Validation score: 0.447153
Iteration 5, loss = 1.42703073
Validation score: 0.445468
Iteration 6, loss = 1.41189643
Validation score: 0.441649
Iteration 5, loss = 1.42703073
Validation score: 0.445468
Iteration 6, loss = 1.41189643
Validation score: 0.441649
Iteration 7, loss = 1.40081894
Validation score: 0.450073
Iteration 8, loss = 1.39218526
Validation score: 0.458160
Iteration 7, loss = 1.40081894
Validation score: 0.450073
Iteration 8, loss = 1.39218526
Validation score: 0.458160
Iteration 9, loss = 1.38299724
Validation score: 0.458722
Iteration 10, loss = 1.37539509
Validation score: 0.441874
Iteration 9, loss = 1.38299724
Validation score: 0.458722

Iteration 10, loss = 1.37539509
Validation score: 0.441874
Iteration 11, loss = 1.36837243
Validation score: 0.450635
Iteration 12, loss = 1.36266357
Validation score: 0.461417
Iteration 11, loss = 1.36837243
Validation score: 0.450635
Iteration 12, loss = 1.36266357
Validation score: 0.461417
Iteration 13, loss = 1.35764530
Validation score: 0.451309
Iteration 14, loss = 1.35227551
Validation score: 0.451870
Iteration 13, loss = 1.35764530
Validation score: 0.451309
Iteration 14, loss = 1.35227551
Validation score: 0.451870
Iteration 15, loss = 1.34790847
Validation score: 0.463552
Iteration 16, loss = 1.34188027
Validation score: 0.456700
Iteration 15, loss = 1.34790847
Validation score: 0.463552
Iteration 16, loss = 1.34188027
Validation score: 0.456700
Iteration 17, loss = 1.33965273
Validation score: 0.470179
Iteration 18, loss = 1.33717135
Validation score: 0.462990
Iteration 17, loss = 1.33965273
Validation score: 0.470179
Iteration 18, loss = 1.33717135
Validation score: 0.462990
Iteration 19, loss = 1.33366985
Validation score: 0.473885
Iteration 20, loss = 1.33063796
Validation score: 0.458946
Iteration 19, loss = 1.33366985
Validation score: 0.473885
Iteration 20, loss = 1.33063796
Validation score: 0.458946
Iteration 21, loss = 1.32808252
Validation score: 0.475458
Iteration 22, loss = 1.32449754
Validation score: 0.473324
Iteration 21, loss = 1.32808252
Validation score: 0.475458
Iteration 22, loss = 1.32449754
Validation score: 0.473324
Iteration 23, loss = 1.32405017
Validation score: 0.458834
Iteration 24, loss = 1.32175343
Validation score: 0.467371
Iteration 23, loss = 1.32405017
Validation score: 0.458834
Iteration 24, loss = 1.32175343
Validation score: 0.467371
Iteration 25, loss = 1.31924153
Validation score: 0.467034
Iteration 26, loss = 1.31772664
Validation score: 0.470628
Iteration 25, loss = 1.31924153
Validation score: 0.467034
Iteration 26, loss = 1.31772664
Validation score: 0.470628
Iteration 27, loss = 1.31647045
Validation score: 0.471976
Iteration 28, loss = 1.31436852

```
Validation score: 0.467146
Iteration 27, loss = 1.31647045
Validation score: 0.471976
Iteration 28, loss = 1.31436852
Validation score: 0.467146
Iteration 29, loss = 1.31278806
Validation score: 0.467146
Iteration 30, loss = 1.31322217
Validation score: 0.463215
Iteration 29, loss = 1.31278806
Validation score: 0.467146
Iteration 30, loss = 1.31322217
Validation score: 0.463215
Iteration 31, loss = 1.30975313
Validation score: 0.465686
Iteration 32, loss = 1.30883224
Validation score: 0.469954
Iteration 31, loss = 1.30975313
Validation score: 0.465686
Iteration 32, loss = 1.30883224
Validation score: 0.469954
Iteration 33, loss = 1.30800027
Validation score: 0.474334
Iteration 34, loss = 1.30708656
Validation score: 0.470179
Iteration 33, loss = 1.30800027
Validation score: 0.474334
Iteration 34, loss = 1.30708656
Validation score: 0.470179
Iteration 35, loss = 1.30669133
Validation score: 0.474222
Iteration 36, loss = 1.30386370
Validation score: 0.473436
Iteration 35, loss = 1.30669133
Validation score: 0.474222
Iteration 36, loss = 1.30386370
Validation score: 0.473436
Iteration 37, loss = 1.30225164
Validation score: 0.459957
Validation score did not improve more than tol=0.000100 for 15 consecutive epochs.
```

✓ Training complete in 5.70 seconds!
Training stopped at iteration: 37

```
=====
NEURAL NETWORK MODEL EVALUATION (Flow-Level)
=====
```

```
Accuracy: 0.4646 (46.46%)
Iteration 37, loss = 1.30225164
Validation score: 0.459957
Validation score did not improve more than tol=0.000100 for 15 consecutive epochs.
```

✓ Training complete in 5.70 seconds!
Training stopped at iteration: 37

```
=====
NEURAL NETWORK MODEL EVALUATION (Flow-Level)
=====
```

```
Accuracy: 0.4646 (46.46%)
```

```
[13]: # Compare all models
print("\n" + "="*60)
print("FINAL MODEL COMPARISON (Flow-Level Classification)")
print("="*60)
print(f"\nDataset Statistics:")
print(f" Total flows: {len(flow_df)}")
```

```

print(f" Training flows: {len(X_train)}")
print(f" Test flows: {len(X_test)}")
print(f" Features per flow: {X_train_processed.shape[1]}")
print(f" Number of classes: {y.unique()}")

print("\nModel Performance:")
print(f" Random Forest Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%) - trained in {train_time}s")
print(f" SVM Accuracy: {accuracy_svm:.4f} ({accuracy_svm*100:.2f}%) - trained in {svm_time}s")
print(f" Neural Network Accuracy: {accuracy_nn:.4f} ({accuracy_nn*100:.2f}%) - trained in {nn_time}s")

print("\nBest Model: ", end="")
best_accuracy = max(accuracy, accuracy_svm, accuracy_nn)
if best_accuracy == accuracy:
    print(f"Random Forest ({accuracy*100:.2f}%)")
elif best_accuracy == accuracy_svm:
    print(f"SVM ({accuracy_svm*100:.2f}%)")
else:
    print(f"Neural Network ({accuracy_nn*100:.2f}%)")

```

```
=====
FINAL MODEL COMPARISON (Flow-Level Classification)
=====
```

Dataset Statistics:

```

Total flows: 111284
Training flows: 89026
Test flows: 22257
Features per flow: 20
Number of classes: 30

```

Model Performance:

```

Random Forest Accuracy: 0.5039 (50.39%) - trained in 1.41s
SVM Accuracy: 0.4198 (41.98%) - trained in 314.05s
Neural Network Accuracy: 0.4646 (46.46%) - trained in 5.70s

```

Best Model: Random Forest (50.39%)

Results

Visualization

Comprehensive visualization of flow-based classification results including model performance, feature importance, and flow statistics.

```
[14]: import matplotlib.pyplot as plt
import seaborn as sns

# Set style
sns.set_style("whitegrid")
sns.set_palette("husl")

fig = plt.figure(figsize=(18, 14))
gs = fig.add_gridspec(3, 3, hspace=0.4, wspace=0.35)

# 1. Model Accuracy Comparison
ax1 = fig.add_subplot(gs[0, :2])
models = ['Random Forest', 'SVM', 'Neural Network']
accuracies = [accuracy * 100, accuracy_svm * 100, accuracy_nn * 100]
colors = ['#2ecc71', '#e74c3c', '#3498db']
```

```

bars = ax1.bar(models, accuracies, color=colors, alpha=0.8, edgecolor='black', linewidth=1)
ax1.set_ylabel('Accuracy (%)', fontsize=13, fontweight='bold')
ax1.set_title('Flow-Level Model Performance Comparison', fontsize=15, fontweight='bold')
ax1.set_ylim([0, 100])
ax1.grid(axis='y', alpha=0.3, linestyle='--')

for bar, acc in zip(bars, accuracies):
    ax1.text(bar.get_x() + bar.get_width()/2., bar.get_height() + 1,
             f'{acc:.2f}%', ha='center', va='bottom', fontweight='bold', fontsize=13)

# 2. Training Time Comparison
ax2 = fig.add_subplot(gs[0, 2])
training_times = [rf_train_time, svm_train_time, nn_train_time]
ax2.bart(models, training_times, color=colors, alpha=0.8, edgecolor='black', linewidth=1)
ax2.set_xlabel('Time (seconds)', fontsize=11, fontweight='bold')
ax2.set_title('Training Time', fontsize=13, fontweight='bold', pad=10)
ax2.grid(axis='x', alpha=0.3, linestyle='--')
for i, t in enumerate(training_times):
    ax2.text(t + 0.5, i, f'{t:.1f}s', ha='center', fontweight='bold', fontsize=10)

# 3. Flow Feature Importance (Random Forest)
ax3 = fig.add_subplot(gs[1, :2])
top_features = feature_importance.head(10)
bars = ax3.bart(range(len(top_features)), top_features['importance'],
                 color='#2ecc71', alpha=0.7, edgecolor='black', linewidth=1.5)
ax3.set_yticks(range(len(top_features)))
ax3.set_yticklabels(top_features['feature'], fontsize=10)
ax3.set_xlabel('Importance Score', fontsize=11, fontweight='bold')
ax3.set_title('Top 10 Flow Features (Random Forest Importance)', fontsize=13, fontweight='bold')
ax3.invert_yaxis()
ax3.grid(axis='x', alpha=0.3, linestyle='--')

for bar, imp in zip(bars, top_features['importance']):
    ax3.text(bar.get_width() + 0.005, bar.get_y() + bar.get_height()/2,
             f'{imp:.4f}', ha='center', va='center', fontweight='bold', fontsize=9)

# 4. Flow Size Distribution
ax4 = fig.add_subplot(gs[1, 2])
flow_sizes = pdf.groupby('flow_key').size()
ax4.hist(flow_sizes.clip(upper=100), bins=50, color="#9b59b6", alpha=0.7, edgecolor='black', linewidth=1)
ax4.set_xlabel('Packets per Flow', fontsize=11, fontweight='bold')
ax4.set_ylabel('Number of Flows', fontsize=11, fontweight='bold')
ax4.set_title('Flow Size Distribution\n(clipped at 100)', fontsize=11, fontweight='bold')
ax4.grid(axis='y', alpha=0.3, linestyle='--')

# 5. Class Distribution
ax5 = fig.add_subplot(gs[2, 0])
class_counts = y_test.value_counts().head(8)
bars = ax5.bart(range(len(class_counts)), class_counts.values, color="#3498db", alpha=0.7)
ax5.set_yticks(range(len(class_counts)))
ax5.set_yticklabels(class_counts.index, fontsize=9)
ax5.set_xlabel('Number of Flows', fontsize=10, fontweight='bold')
ax5.set_title('Top 8 Application Classes\n(n(Test Set))', fontsize=11, fontweight='bold')
ax5.invert_yaxis()
ax5.grid(axis='x', alpha=0.3, linestyle='--')

# 6. Packet Size Distribution by Flow
ax6 = fig.add_subplot(gs[2, 1])
flow_df['pkt_len_mean'].hist(bins=50, color="#e67e22", alpha=0.7)
ax6.set_xlabel('Mean Packet Size (bytes)', fontsize=10, fontweight='bold')
ax6.set_ylabel('Number of Flows', fontsize=10, fontweight='bold')
ax6.set_title('Flow Mean Packet Size\nDistribution', fontsize=11, fontweight='bold')

```

```

ax6.grid(axis='y', alpha=0.3, linestyle='--')

# 7. Summary Statistics
ax7 = fig.add_subplot(gs[2, 2])
ax7.axis('tight')
ax7.axis('off')
ax7.set_title('Flow-Based Analysis Summary', fontsize=11, fontweight='bold', pad=2

summary_data = [
    ['Metric', 'Value', ''],
    ['Total Packets', f'{len(pdf):,}', ''],
    ['Total Flows', f'{len(flow_df):,}', ''],
    ['Train Flows', f'{len(X_train):,}', ''],
    ['Test Flows', f'{len(X_test):,}', ''],
    ['Flow Features', f'{X_train_processed.shape[1]}', ''],
    ['Classes', f'{y.unique()}', ''],
    ['', '', ''],
    ['Model', 'Accuracy', 'Rank'],
    ['Random Forest', f'{accuracy*100:.2f}%', '1st' if accuracy == best_accuracy else ''],
    ['Neural Network', f'{accuracy_nn*100:.2f}%', '1st' if accuracy_nn == best_accuracy else ''],
    ['SVM', f'{accuracy_svm*100:.2f}%', '1st' if accuracy_svm == best_accuracy else ''],
]
]

table = ax7.table(cellText=summary_data, cellLoc='left', loc='upper center',
                  colWidths=[0.4, 0.3, 0.2], bbox=[0, 0, 1, 0.85])
table.auto_set_font_size(False)
table.set_fontsize(8)
table.scale(1, 1.8)

# Style header rows
for i in [0, 8]:
    for j in range(3):
        table[(i, j)].set_facecolor('#34495e')
        table[(i, j)].set_text_props(weight='bold', color='white')

fig.suptitle('Application Identification in Network Traffic - Flow-Based Analysis',
             fontsize=16, fontweight='bold', y=0.98)

plt.tight_layout(rect=[0, 0, 1, 0.97])
plt.savefig('flow_based_results_visualization.png', dpi=300, bbox_inches='tight')
plt.show()

print("✓ Visualization saved as 'flow_based_results_visualization.png'")
```

Application Identification in Network Traffic - Flow-Based Analysis Results



✓ Visualization saved as 'flow_based_results_visualization.png'

Summary of Findings

This project explored three machine learning approaches for identifying applications in encrypted network traffic using **flow-level features**. Unlike packet-level analysis, flow-based classification aggregates related packets (using the 5-tuple: src_ip, dst_ip, src_port, dst_port, protocol) and extracts statistical features that capture traffic behavior patterns.

Dataset Statistics

- **Total packets:** 529,019
- **Total flows extracted:** 111,284
- **Average packets per flow:** 4.75 (median: 2.0)
- **Flow size range:** 1 to 875 packets
- **Features per flow:** 20
- **Application classes:** 30

The dataset includes traffic from various applications including Hangouts audio (126K packets), Facebook audio (122K packets), Skype file transfers (85K packets), and others.

Flow Definition and Extraction

A **network flow** was defined as a bidirectional sequence of packets sharing the same 5-tuple. For each flow, we computed 20 statistical features:

- **Packet statistics:** total packets, total bytes, packets per second
- **Size statistics:** mean, std, min, max, median packet lengths
- **Timing statistics:** flow duration, inter-arrival time (IAT) mean/std/min/max
- **Packet size distribution:** counts of small (<100B), medium (100-500B), and large (>500B) packets
- **Protocol features:** protocol type, port numbers

Model Performance

The three models achieved the following results on flow-level classification (89,026 training flows, 22,257 test flows):

Model	Accuracy	Training Time
Random Forest	50.39%	1.41s
Neural Network (MLP)	46.46%	5.70s
Support Vector Machine	41.98%	314.05s

Best Model: Random Forest with 50.39% accuracy across 30 application classes.

Top 10 Most Important Flow Features (Random Forest)

Rank	Feature	Importance
1	max_port	0.1753
2	bytes_per_sec	0.1036
3	iat_max	0.0925
4	flow_duration	0.0702
5	packets_per_sec	0.0623
6	iat_std	0.0573
7	pkt_len_std	0.0547
8	pkt_len_median	0.0523
9	iat_mean	0.0522
10	total_bytes	0.0490

The most discriminative features are port numbers (likely indicating server-side services), flow rate metrics, and inter-arrival time patterns.

Key Observations

- **Flow-level vs Packet-level:** Flow-based analysis provides a more natural unit for application classification. Aggregating 529K packets into 111K flows reduces complexity while capturing meaningful patterns.
- **Feature Importance:** Port numbers and timing features (IAT, flow duration) are the most discriminative, followed by packet size statistics. This suggests applications have distinctive temporal behavior patterns.
- **Class Imbalance:** The 30-class classification is challenging, with some applications having significantly more samples than others, contributing to the ~50% accuracy ceiling.
- **Model Trade-offs:** Random Forest provides the best accuracy with minimal training time. SVM required 314 seconds but achieved lower accuracy, suggesting linear boundaries are insufficient for this

problem.

Practical Implications

Flow-based application identification at ~50% accuracy across 30 classes demonstrates:

- **Feasibility:** Machine learning can distinguish applications from encrypted traffic metadata alone
- **Network Management:** Even partial accuracy enables traffic prioritization and QoS policies
- **Security Monitoring:** Unusual flow patterns can flag potential threats
- **Privacy Preservation:** No payload inspection required

Conclusion

This project successfully demonstrated flow-level machine learning for application identification in encrypted network traffic. By defining flows using the 5-tuple and extracting 20 statistical features, we achieved **50.39% accuracy** with Random Forest across **30 application classes** on **111,284 flows**.

The flow-based approach directly addresses the feedback about defining and extracting flows from PCAPNG data, providing a rigorous and industry-standard methodology for traffic analysis.

Future work could improve accuracy through:

- Deep learning on raw packet sequences within flows (RNN/LSTM)
- Addressing class imbalance with oversampling/undersampling
- Adding more sophisticated features (e.g., packet direction, burst patterns)
- Ensemble methods combining multiple classifiers
- Transfer learning across different network environments